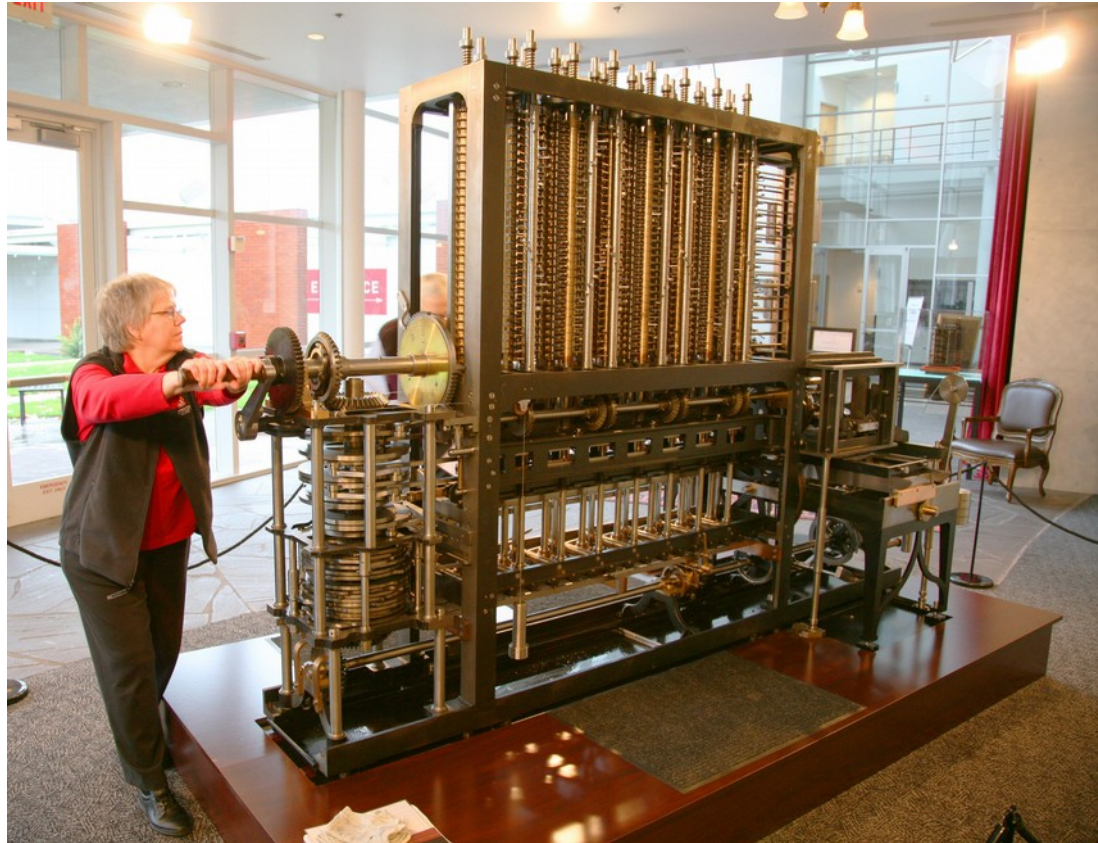# PPAR: CUDA basics

Sylvain Collange
Inria Rennes – Bretagne Atlantique
sylvain.collange@inria.fr

PPAR 2017

This lecture: CUDA programming

- We have seen some GPU architecture

- Now how to program it ?

Outline

- GPU programming environments

- CUDA host side

- CUDA device side: threads, blocks, grids

- Expressing parallelism

  - Vector add example

- Managing communications

  - Parallel reduction example

- Re-using data

  - Matrix multiplication example

GPU development environments

For general-purpose programming (not graphics)

- Multiple toolkits
  - NVIDIA CUDA
  - Khronos OpenCL
  - Vulkan Compute
  - Microsoft DirectCompute
  - Google RenderScript
- Mostly syntactical variations
  - Underlying principles are the same
- In this course, focus on NVIDIA CUDA

Higher-level programming

- Directive-based
  - OpenACC
  - OpenMP 4.x

- Language extensions / libraries
  - Microsoft C++ AMP
  - Intel Cilk+
  - NVIDIA Thrust, CUB

- Languages
  - Intel ISPC

...

- Most corporations agree we need common standards...
  - But only if their own product becomes the standard!

Outline

- GPU programming environments

- CUDA host side

- Parallel programming models: BSP, multi-BSP

- CUDA device side: threads, blocks, grids

- Expressing parallelism

  - Vector add example

- Managing communications

  - Parallel reduction example

- Re-using data

  - Matrix multiplication example

Hello World in CUDA

- CPU "host" code + GPU "device" code

hello.cu:

```
__global__ void hello() {
}
```
} Device code

```
int main() {
    hello<<<1,1>>>();
    printf("Hello World!\n");
    return 0;
}
```
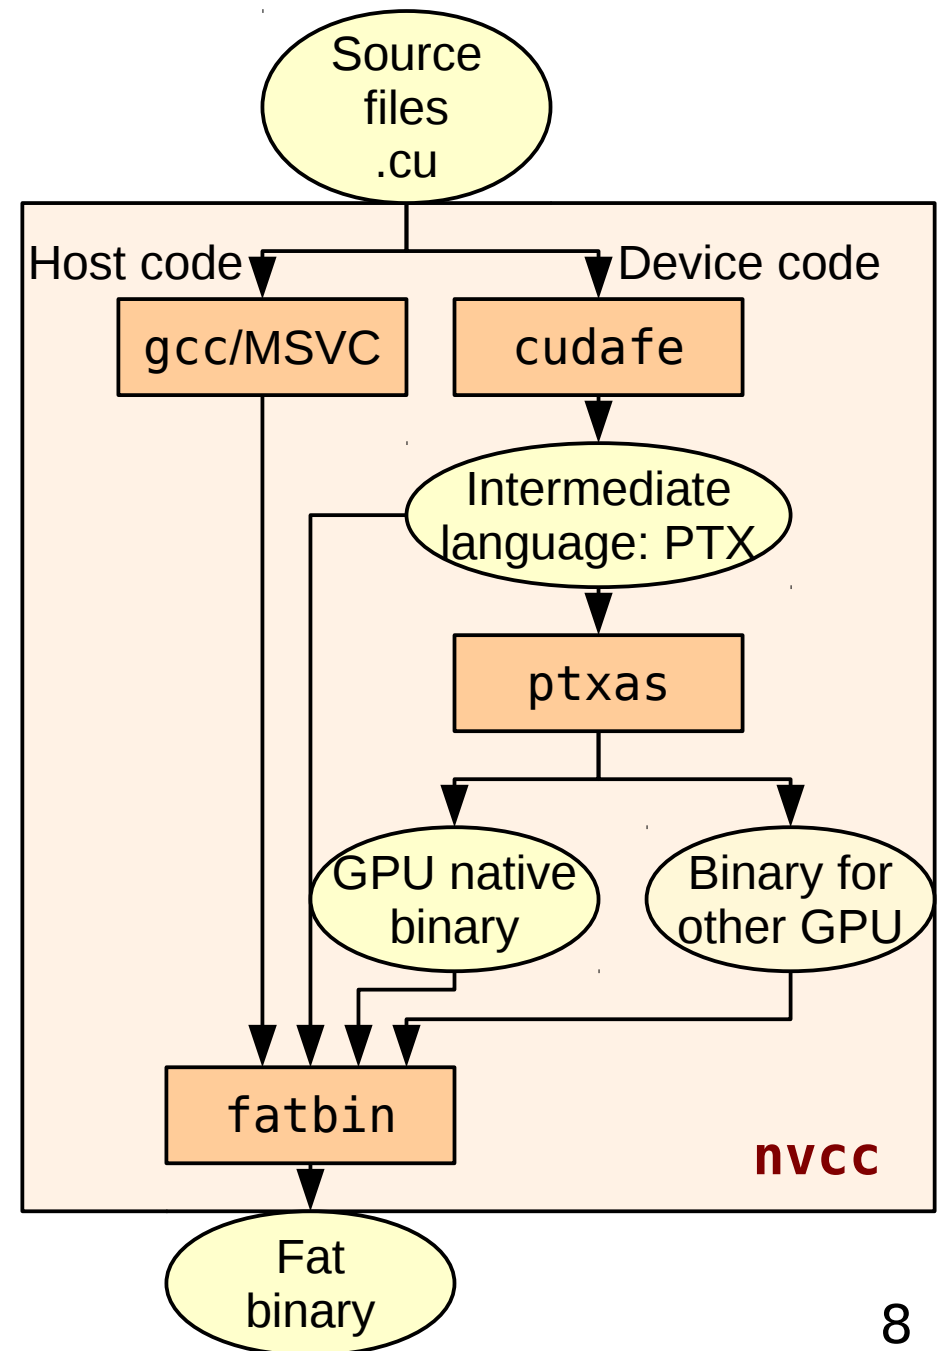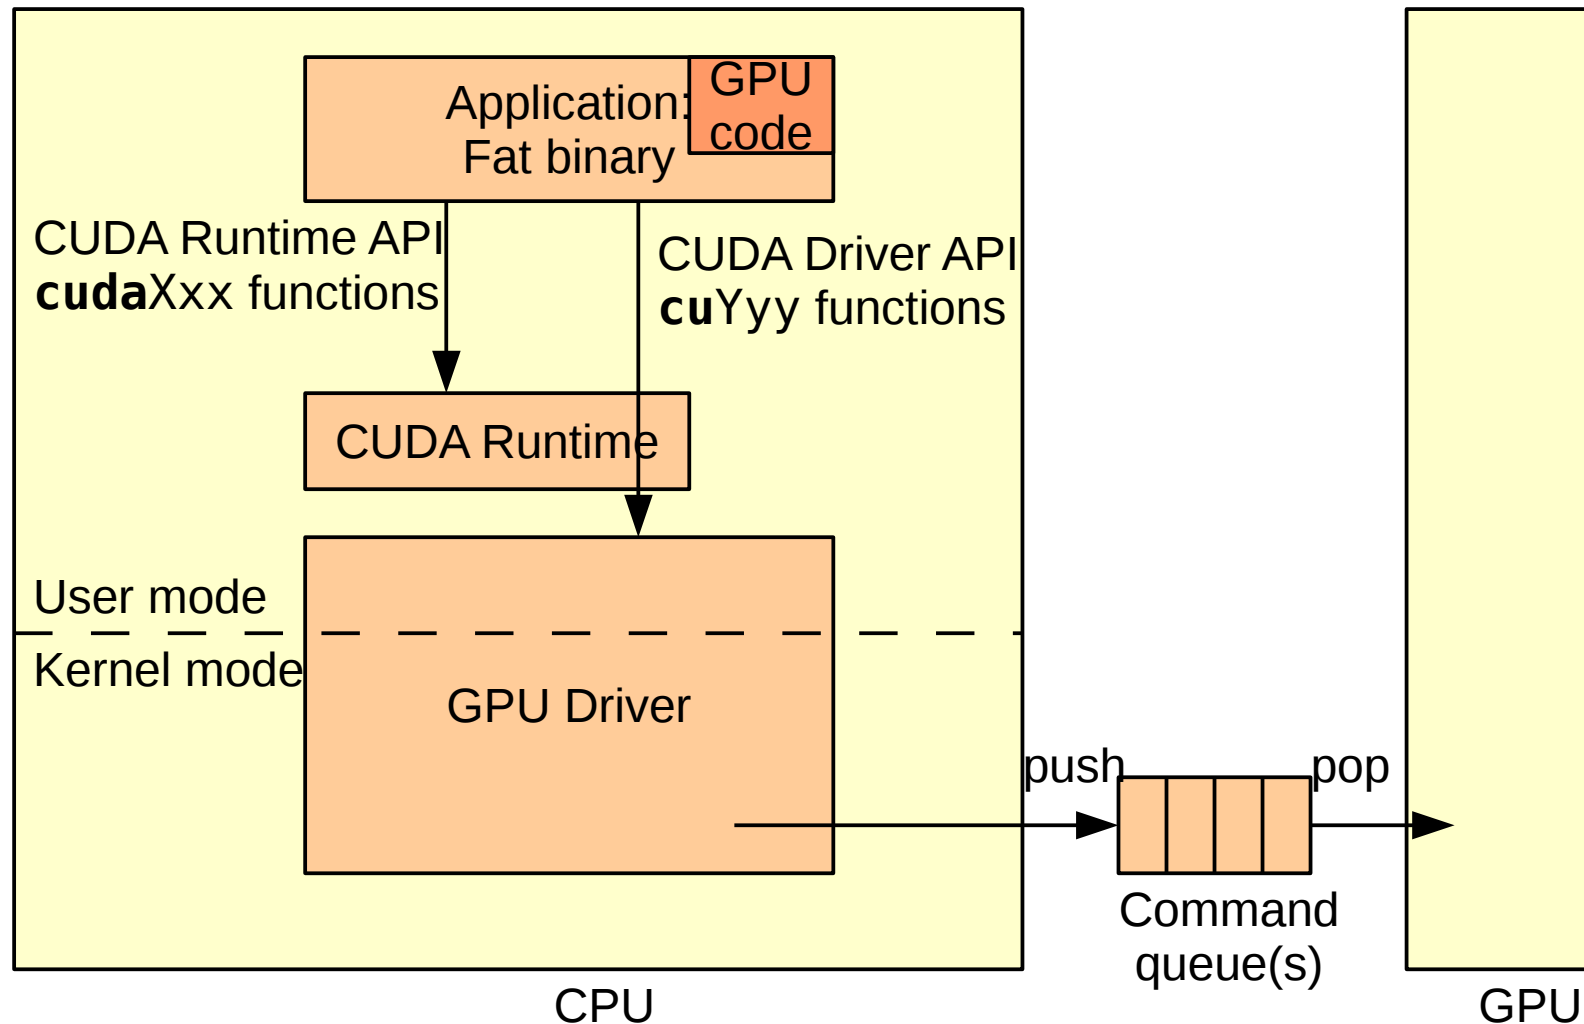} Host code

Compiling a CUDA program

- Executable contains both host and device code
  - Device code in PTX and/or native
  - PTX can be recompiled on the fly
    (e.g. old program on new GPU)
- NVIDIA's compiler driver takes care of the process:

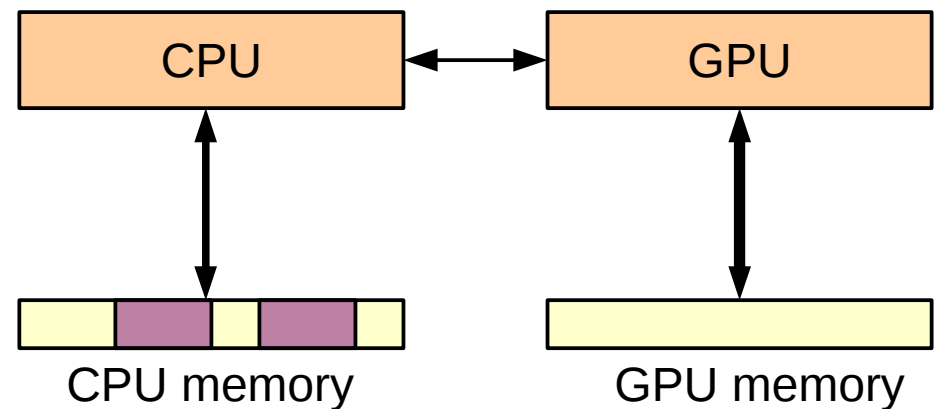 **nvcc** -o hello hello.cu

8

Control flow

- Program running on CPUs

- Submit work to the GPU through the GPU driver
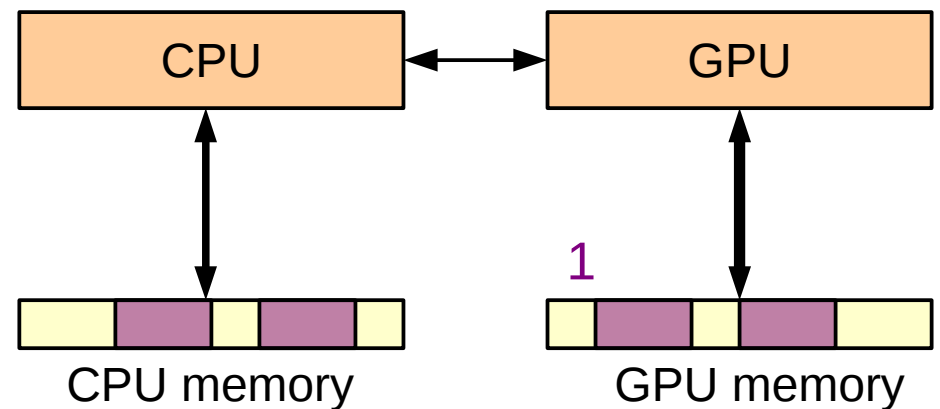
- Commands execute asynchronously

Data flow

- Main program runs on the host
  - Manages memory transfers
  - Initiate work on GPU
- Typical flow

Data flow

- Main program runs on the host
  - Manages memory transfers
  - Initiate work on GPU
- Typical flow
  - **1.** Allocate GPU memory

11

Data flow

- Main program runs on the host
    - Manages memory transfers
    - Initiate work on GPU
- Typical flow
    - **1.** Allocate GPU memory
    - **2.** Copy inputs from CPU mem to GPU memory

Data flow

- Main program runs on the host

  - Manages memory transfers

  - Initiate work on GPU

- Typical flow

  - **1.** Allocate GPU memory

  - **2.** Copy inputs from CPU mem to GPU memory
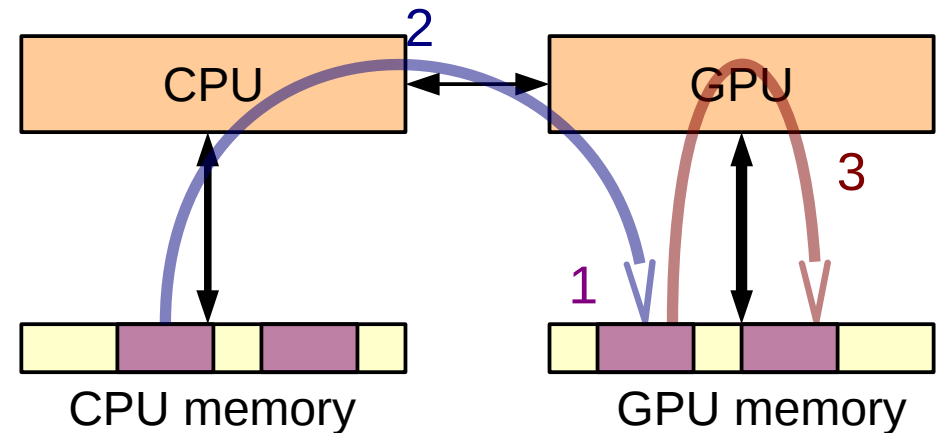
  - **3.** Run computation on GPU

Data flow

- Main program runs on the host
  - Manages memory transfers
  - Initiate work on GPU
- Typical flow
  - **1.** Allocate GPU memory
  - **2.** Copy inputs from CPU mem to GPU memory
  - **3.** Run computation on GPU
  - **4.** Copy back results to CPU memory

Example: a + b

- Our Hello World example did not involve the GPU
- Let's add up 2 numbers on the GPU
- Start from host code

```
int main()
{
    float ab[2] = {1515, 149};   // Inputs

    float c[1];  // Output
    // c[0] = ab[0] + ab[1];
    printf("c = %f\n", c[0]);
}
```

vectorAdd example: cuda/samples/0_Simple/vectorAdd

Step 1: allocate GPU memory

```c
int main()
{
    float ab[2] = {1515, 149}, c[1];  // Inputs, in host mem

    // Allocate GPU memory
    float *d_AB, *d_C;
    cudaMalloc((void **)&d_AB, 2*sizeof(float));
    cudaMalloc((void **)&d_C, sizeof(float));
```

Passing a pointer to the
pointer to be overwritten

```c
    // Free GPU memory
    cudaFree(d_AB);
    cudaFree(d_C);
}
```

- Allocate space for a, b and c in GPU memory

- At the end, free memory

16

Step 2, 4: copy data to/from GPU memory

```c
int main()
{
    float ab[2] = {1515, 149}, c[1];    // Inputs/outputs, CPU mem

    // Allocate GPU memory
    float *d_AB, *d_C;
    cudaMalloc((void **)&d_AB, 2*sizeof(float));
    cudaMalloc((void **)&d_C, sizeof(float));


    // Copy from CPU mem to GPU mem
    cudaMemcpy(d_AB, ab, 2*sizeof(float), cudaMemcpyHostToDevice);



    // Copy results back to CPU mem
    cudaMemcpy(c, d_C, sizeof(float), cudaMemcpyDeviceToHost);
    printf("c = %f\n", c[0]);


    // Free GPU memory
    cudaFree(d_AB);
    cudaFree(d_C);
}
```

Step 3: launch kernel

```c
__global__ void addOnGPU(float * ab, float * c)
{
    c[0] = ab[0] + ab[1];
}


int main()
{
    float ab[] = {1515, 159};    // Inputs, CPU mem

    // Allocate GPU memory
    float *d_AB, *d_C;
    cudaMalloc((void **)&d_AB, 2*sizeof(float));
    cudaMalloc((void **)&d_C, sizeof(float));
    // Copy from CPU mem to GPU mem
    cudaMemcpy(d_AB, ab, 2*sizeof(float), cudaMemcpyHostToDevice);



    // Launch computation on GPU
    addOnGPU<<<1, 1>>>(d_AB, d_C);

    float c[1];     // Result on GPU

    // Copy results back to CPU mem
    cudaMemcpy(c, d_C, sizeof(float), cudaMemcpyDeviceToHost);
    printf("c = %f\n", c[0]);
    // Free GPU memory
    cudaFree(d_AB);
    cudaFree(d_C);
}
```

- Kernel is a function prefixed by `__global__`
    - Runs on GPU
- Invoked from CPU code with `<<<>>>` syntax

Note: we could have passed a and b directly as kernel parameters

What is inside the <<<>>>?

18

Outline

- GPU programming environments

- CUDA host side

- CUDA device side: threads, blocks, grids

- Expressing parallelism

  - Vector add example

- Managing communications

  - Parallel reduction example

- Re-using data

 - Matrix multiplication example

Workload: logical organization

- A kernel is launch on a grid: `my_kernel<<<blocks, threads>>>(...)`
- Two nested levels
  - Blocks
  - Threads
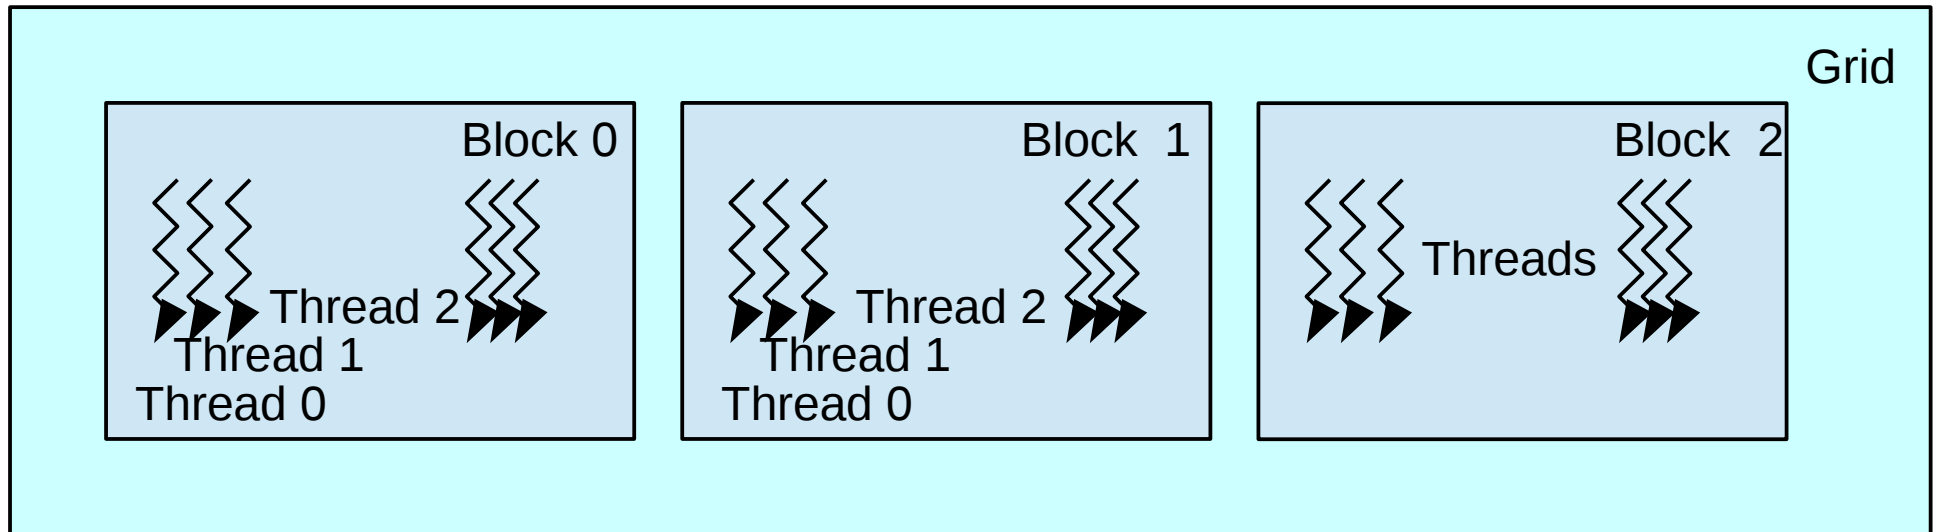
Outer level: grid of blocks

- Blocks also named Concurrent Thread Arrays (CTAs)
- **No communication** between blocks of the same grid
- Practically **unlimited number** of blocks / grid

Grid 1

Block 1

Block 2

...

Block 3

Inner level: threads

- Blocks contain threads

- All threads in a block

  - Run on the same SM:
    they can **communicate**

  - Run in parallel:
    they can **synchronize**

- Constraints on
  number of threads / block

  - Maximum:
    512 to 1024 depending on arch

  - Recommended:
    at least 64 threads for good
    performance

  - Recommended:
    multiple of the warp size (32)

Block

Threads

Barrier

23

Multi-BSP and CUDA

Minor difference: BSP is based on message passing, CUDA on shared memory

24
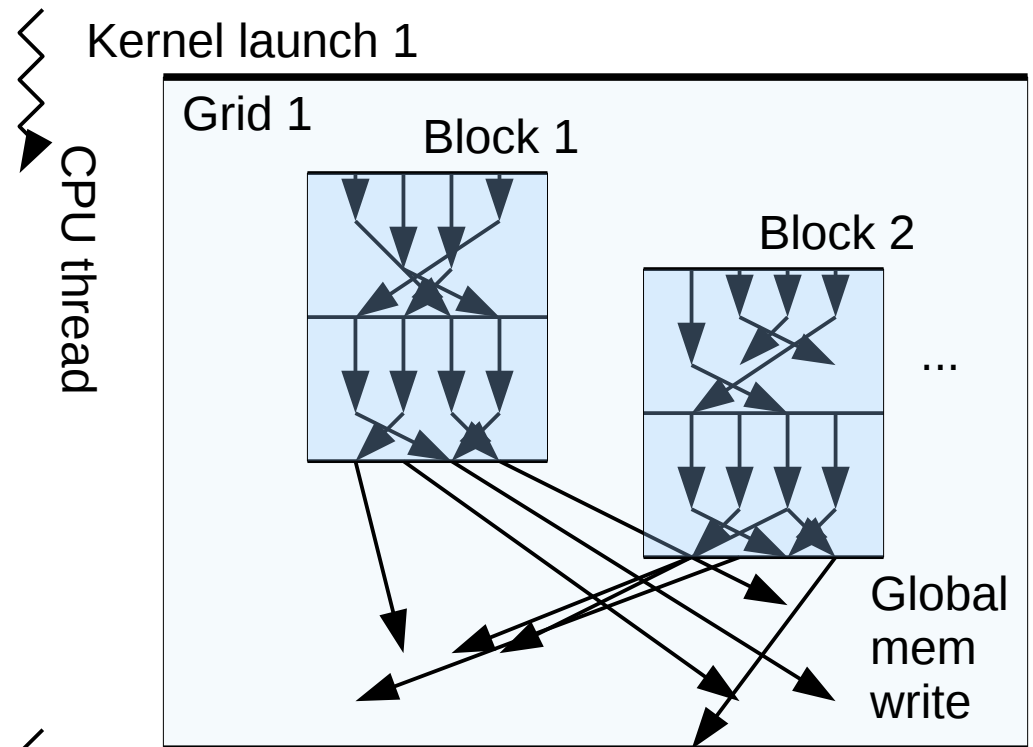
GPU physical organization
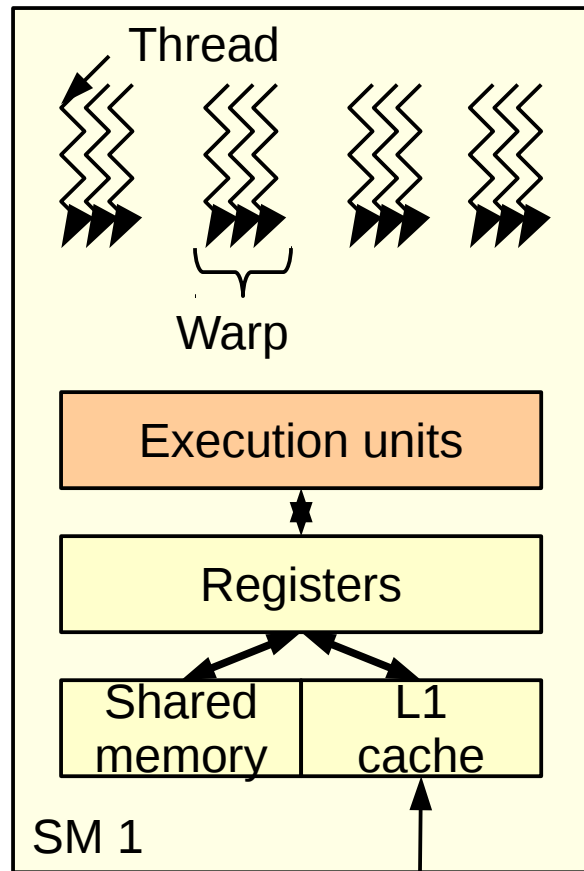
Mapping blocks to hardware resources
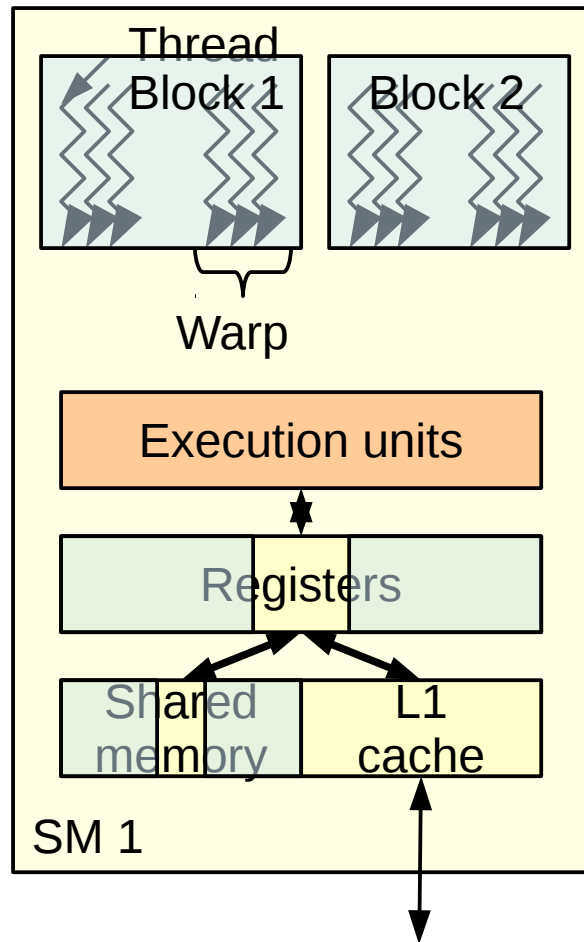
- SM resources are partitioned across blocks

Block scheduling

- Blocks may
  - Run serially or in parallel
  - Run on the same or different SM
  - Run in order or out of order
- Should not assume anything on execution order of blocks

Blocks

0	1	2	3
4	5	6	7
8	9	10	11

SM 1   SM 2

Block scheduling

- Blocks may
  - Run serially or in parallel
  - Run on the same or different SM
  - Run in order or out of order
- Should not assume anything on execution order of blocks

Blocks

5	6	7
8	9	10	11

SM 1

| 0 | 2 |
| 4 |

SM 2

| 1 | 3 |

Block scheduling

- Blocks may
  - Run serially or in parallel
  - Run on the same or different SM
  - Run in order or out of order
- Should not assume anything on execution order of blocks

Blocks

| 8 | 9 | 10 | 11 |

| 0 | 5 |
| 4 | |

SM 1

| 3 |
| 6 | 7 |

SM 2

| 1 | 2 |

Block scheduling

- Blocks may
  - Run serially or in parallel
  - Run on the same or different SM
  - Run in order or out of order
- Should not assume anything on execution order of blocks

Blocks

SM 1   SM 2

Outline

- GPU programming environments

- CUDA host side

- CUDA device side: threads, blocks, grids

- **Expressing parallelism**

  - Vector add example

- **Managing communications**

  - Parallel reduction example

- **Re-using data**

  - Matrix multiplication example

Example: vector addition

- Addition example: only 1 thread

  - Now let's run a parallel computation

- Start with multiple blocks, 1 thread/block

  - Independent computations in each block

- No communication/synchronization needed

Host code: initialization

- A and B are now arrays: just change allocation size

```c
int main()
{
    int numElements = 50000;
    size_t size = numElements * sizeof(float);

    float *h_A = (float *)malloc(size);
    float *h_B = (float *)malloc(size);
    float *h_C = (float *)malloc(size);
    Initialize(h_A, h_B);

    // Allocate device memory
    float *d_A, *d_B, *d_C;
    cudaMalloc((void **)&d_A, size);
    cudaMalloc((void **)&d_B, size);
    cudaMalloc((void **)&d_C, size);

    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
    ...
}
```

33

Host code: kernel and kernel launch

```
__global__ void vectorAdd2(float *A, float *B, float *C)
{
    int i = blockIdx.x;

    C[i] = A[i] + B[i];
}
```

- Launch n blocks of 1 thread each (for now)

```
int blocks = numElements;
vectorAdd2<<<blocks, 1>>>(d_A, d_B, d_C);
```

Device code

```
__global__ void vectorAdd2(float *A, float *B, float *C)
{
    int i = blockIdx.x;          Built-in CUDA variable:
                                 in device code only

    C[i] = A[i] + B[i];
}
```

- Block number *i* processes element *i*
- Grid of blocks may have up to 3 dimensions (`blockIdx.x`, `blockIdx.y`, `blockIdx.z`)
    - For programmer convenience: no effect on scheduling

Multiple blocks, multiple threads/block

Fixed number of threads / block: here 64

- Host code

Not necessarily multiple of block size!

```
int threads = 64;
int blocks = (numElements + threads - 1) / threads; // Round up

vectorAdd3<<<blocks, threads>>>(d_A, d_B, d_C, numElements);
```

- Device code

```
__global__ void vectorAdd3(const float *A, const float *B, float *C,
    int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;

    if(i < n) {
        C[i] = A[i] + B[i];
    }
}
```

Global index

Last block may have less work to do

Thread block may also have up to 3 dimensions: `threadIdx.{x,y,z}`   36

Outline

- GPU programming environments

- CUDA host side

- CUDA device side: threads, blocks, grids

- Expressing parallelism

  - Vector add example

- **Managing communications**

  - Parallel reduction example

- **Re-using data**

  - Matrix multiplication example

Barriers

- Threads can synchronize inside one block

  - Wait until all threads in the block have reached the barrier

- In C for CUDA:

  ```
  __syncthreads();
  ```

- Needs to be called at the same place
  for all threads of the block

```
if(tid < 5) {          if(a[0] == 17) {        if(tid < 5) {
    ...                    __syncthreads();        __syncthreads();
}                      }                       }
else {                 else {                  else {
    ...                    __syncthreads();        __syncthreads();
}                      }                       }
__syncthreads();
```

Same condition
for all threads in the block

Correct   **Correct**   **Wrong**

38

Shared memory

- Fast, software-managed memory
  - Faster than global memory

- Valid only inside one block
  - Each block sees its own copy

- Used to exchange data between threads

- Concurrent writes:
  one thread wins, but we do not know which one

Thread communication: common pattern

- Each thread writes to its own location
  - No write conflict
- Barrier
  - Wait until all threads have written
- Read data from other threads

Compute

Write to smem[tid]

Barrier

Read from smem[f(tid)]

Barrier

Compute
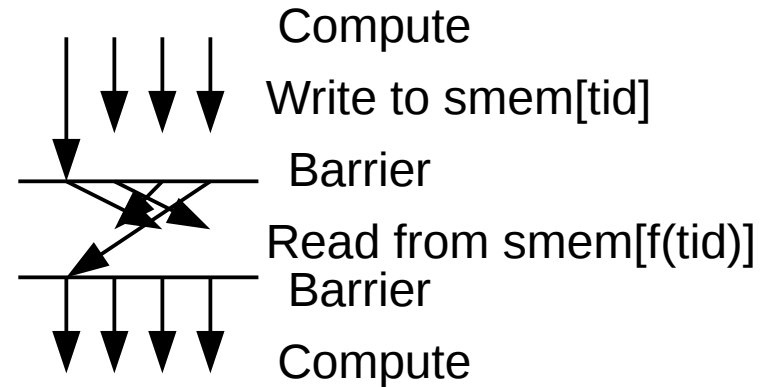
Example: parallel reduction

- Algorithm for 2-level multi-BSP model

Reduction in CUDA: level 1

```
__global__ void reduce1(float *g_idata, float *g_odata, unsigned int n)
{
    extern __shared__ float sdata[];                    Dynamic shared memory allocation:
                                                        will specify size later
    unsigned int tid = threadIdx.x;
    unsigned int i = blockIdx.x * blockDim.x + threadIdx.x;

    // Load from global to shared mem
    sdata[tid] = (i < n) ? g_idata[i] : 0;
    __syncthreads();

    for(unsigned int s = 1; s < blockDim.x; s *= 2) {
        int index = 2 * s * tid;

        if(index < blockDim.x) {
            sdata[index] += sdata[index + s];
        }
        __syncthreads();
    }

    // Write result for this block to global mem
    if (tid == 0) g_odata[blockIdx.x] = sdata[0];
}
```

cuda/samples/6_Advanced/reduction

Reduction: host code

```
int smemSize = threads * sizeof(float);
reduce1<<<blocks, threads, smemSize>>>(d_idata, d_odata, size);
```

Optional parameter:
Size of dynamic shared memory per block

- Level 2: run reduction kernel again,
  until we have 1 block left

- By the way, is our reduction operator associative?

# A word on floating-point

- Parallel reduction requires the operator to be associative

- Is addition associative?

  - On reals: yes

  - On floating-point numbers: no
    With 4 decimal digits:
    (1.234 + 123.4) - 123.4 = 124.6 - 123.4 = **1.200**
    1.234 + (123.4 - 123.4) = 1.234 + 0 = **1.234**

- Consequence: different result depending on thread count

Outline

- GPU programming environments

- CUDA host side

- Parallel programming models: BSP, multi-BSP

- CUDA device side: threads, blocks, grids

- Expressing parallelism

  - Vector add example

- Managing communications

  - Parallel reduction example

- **Re-using data**

  - **Matrix multiplication example**

Arithmetic intensity

Arithmetic throughput

4612 Gflops/s

Arithmetic intensity

114 flops / word

- Example from first lecture
  - NVIDIA GTX 980 needs ≥114 flops / word to reach peak performance
- How to reach enough arithmetic intensity?
  - Need to **reuse** values loaded from memory

Classic example: matrix multiplication

- Naive algorithm

```
for i = 0 to n-1
    for j = 0 to n-1
        for k = 0 to n-1
            C[i,j]+=A[i,k]*B[k,j]
```


- Arithmetic intensity: 1:1   :(

Reusing inputs

- Move loop on *k* up

```
for k = 0 to n-1
    for i = 0 to n-1
        for j = 0 to n-1
            C[i,j]+=A[i,k]*B[k,j]
```


- Enable data reuse on inputs A and B
- But no more reuse on matrix C!

With tiling

- Block loops on i and j

```
for i = 0 to n-1 step 16
    for j = 0 to n-1 step 16
        for k = 0 to n-1
            for i2 = i to i+15
                for j2 = j to j+15
                    C[i2,j2]+=A[i2,k]*B[k,j2]
```

- For one block: product between horizontal panel of A and vertical panel of B

A

B

C

i,j

Constant size

With more tiling

- Block loop on k

```
for i = 0 to n-1 step 16
    for j = 0 to n-1 step 16
        for k = 0 to n-1 step 16
            for k2 = k to k+15
                for i2 = i to i+15
                    for j2 = j to j+15
                        C[i2,j2]+=A[i2,k]*B[k,j2]
```

A
B
C

k

k

Constant size

Pre-loading data

```
for i = 0 to n-1 step 16
    for j = 0 to n-1 step 16
        c = {0}
        for k = 0 to n-1 step 16
            a = A[i..i+15,k..k+15]
            b = B[k..k+15,j..j+15]
```
⎤ Load submatrices a and b

```
            for k2 = 0 to 15
                for i2 = 0 to 15
                    for j2 = 0 to 15
                        c[i2,j2]+=a[i2,k2]*b[k2,j2]
```
⎫ Multiply submatrices
⎬ c = a × b

```
        C[i..i+15,j..j+15] = c
```
⎤ Store submatrix c

A   B   C

Arithmetic intensity?

51

Breaking into two levels

- Run loops on i, j, i2, j2 in parallel

```
for // i = 0 to n-1 step 16          Level 2:
    for // j = 0 to n-1 step 16      Blocks
        c = {0}
        for k = 0 to n-1 step 16
            a = A[i..i+15,k..k+15]              Level 1:
            b = B[k..k+15,j..j+15]              Threads

            for k2 = 0 to 15
                for // i2 = 0 to 15
                    for // j2 = 0 to 15
                        c[i2,j2]+=a[i2,k2]*b[k2,j2]

        C[i..i+15,j..j+15] = c
```
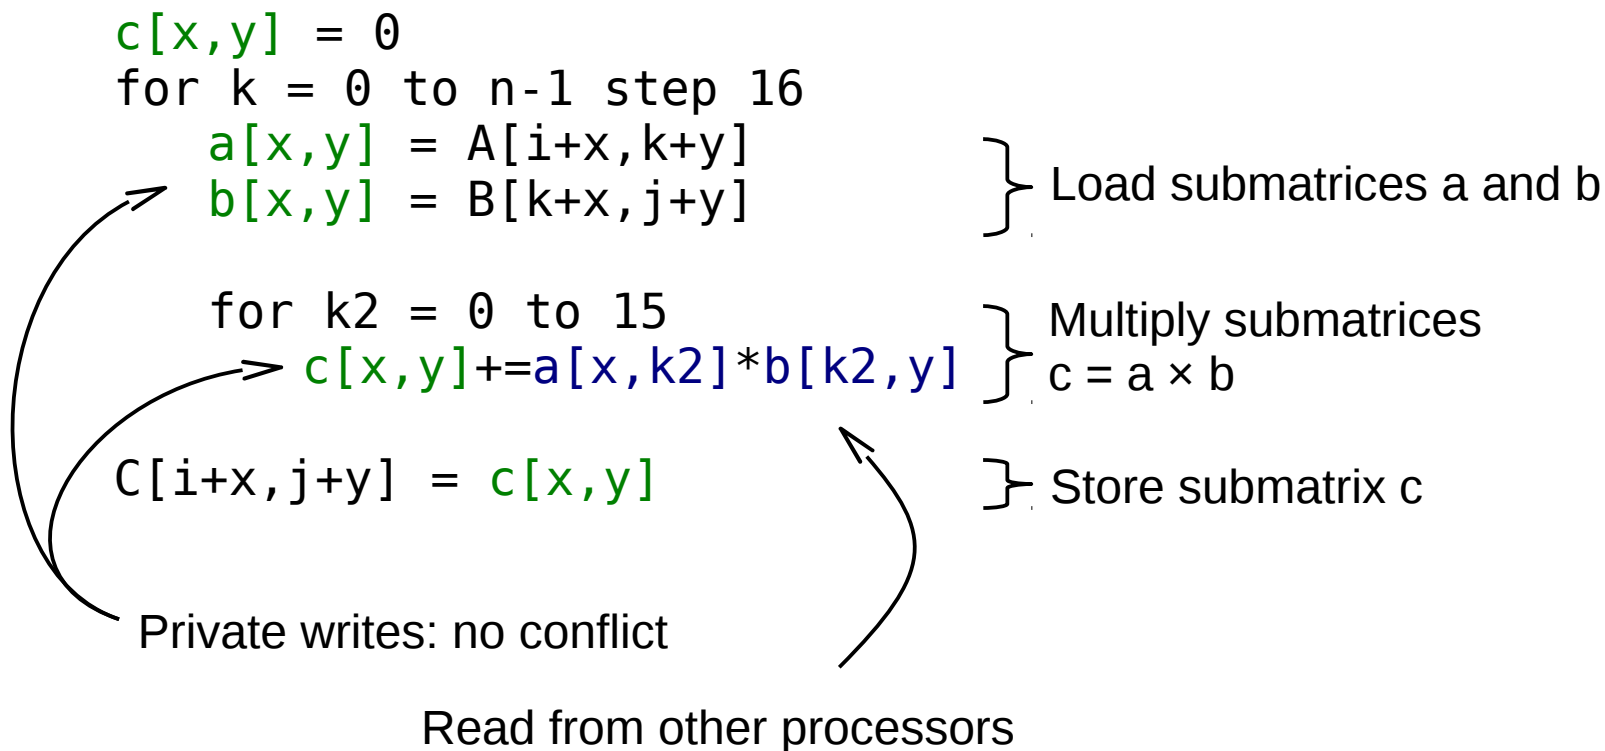
- Let's focus on threads

Level 1: SIMD (PRAM-style) version

- Each processor has ID (x,y)
  - Loops on i2, j2 are implicit

```
c[x,y] = 0
for k = 0 to n-1 step 16
    a[x,y] = A[i+x,k+y]
    b[x,y] = B[k+x,j+y]          }  Load submatrices a and b

    for k2 = 0 to 15             ]  Multiply submatrices
        c[x,y]+=a[x,k2]*b[k2,y]  }  c = a × b

C[i+x,j+y] = c[x,y]              }  Store submatrix c
```

Private writes: no conflict

Read from other processors

- How to translate to SPMD (BSP-style) ?

SPMD version

- Place synchronization barriers

```
c[x,y] = 0
for k = 0 to n-1 step 16
    a[x,y] = A[i+x,k+y]
    b[x,y] = B[k+x,j+y]
    Barrier
    for k2 = 0 to 15
        c[x,y]+=a[x,k2]*b[k2,y]
    Barrier
C[i+x,j+y] = c[x,y]
```


- Why do we need the second barrier ?

Data allocation

- 3 memory spaces: Global, Shared, Local

  - Where should we put: `A, B, C, a, b, c` ?

```
c[x,y] = 0
for k = 0 to n-1 step 16
    a[x,y] = A[i+x,k+y]
    b[x,y] = B[k+x,j+y]
    Barrier
    for k2 = 0 to 15
        c[x,y] += a[x,k2]*b[k2,y]
    Barrier
C[i+x,j+y] = c[x,y]
```

Data allocation

- Memory spaces: Global, Shared, Local
  - As local as possible

```
c = 0
for k = 0 to n-1 step 16
    a[x,y] = A[i+x,k+y]
    b[x,y] = B[k+x,j+y]
    Barrier
    for k2 = 0 to 15
        c += a[x,k2]*b[k2,y]
    Barrier
C[i+x,j+y] = c
```

Local: private to each thread (indices are implicit)

Global: shared between blocks / inputs and outputs

Shared: shared between threads, private to block

CUDA version

- Straightforward translation

```
float Csub = 0;

for(int a = aBegin, b = bBegin;
     a <= aEnd;
     a += aStep, b += bStep) {
     __shared__ float As[BLOCK_SIZE][BLOCK_SIZE];
     __shared__ float Bs[BLOCK_SIZE][BLOCK_SIZE];

    As[ty][tx] = A[a + wA * ty + tx];
    Bs[ty][tx] = B[b + wB * ty + tx];

    __syncthreads();
    for(int k = 0; k < BLOCK_SIZE; ++k)
    {
        Csub += As[ty][k] * Bs[k][tx];
    }
    __syncthreads();
}

int c = wB * BLOCK_SIZE * by + BLOCK_SIZE * bx;
C[c + wB * ty + tx] = Csub;
```

Precomputed base addresses

Declare shared memory

Linearized arrays

matrixMul example: cuda/samples/0_Simple/matrixMul

57

Local memory

- Registers are fast but
  - Limited in size
  - Not addressable

- Local memory used for
  - Local variables that do not fit in registers (*register spilling*)
  - Local arrays accessed with indirection

    ```
    int a[17];
    b = a[i];
    ```

- **Warning**: local is a misnomer!
  - Physically, local memory usually goes off-chip

Device functions

- Kernel can call functions

- Need to be marked for GPU compilation

```
__device__ int foo(int i) {
}
```

- A function can be compiled for both host and device

```
__host__ __device__ int bar(int i) {
}
```

- Device functions can call device functions

  - Older GPUs do not support recursion

Recap

- Memory management:
  Host code and memory / Device code and memory

- Writing GPU Kernels

- Dimensions of parallelism: grids, blocks, threads

- Memory spaces: global, local, shared memory

- Next time:

  - Optimizing memory access

  - Advanced features and optimization techniques